

Project Report: Multilingual NCERT Doubt Solver

1. Project Overview

This project is a **local, offline, RAG-based (Retrieval-Augmented Generation)** AI system designed to answer student questions based **strictly** on NCERT textbooks. It supports multilingual queries, provides precise citations, and formats answers with clear "word equations" and structured explanations suitable for school students (Grades 5-10).

Key Features:

- **Strict RAG:** Answers are grounded purely in the provided PDF textbooks. Hallucination is minimized.
- **Multilingual Support:** Can process and answer queries in multiple languages (English/Hindi) using cross-lingual embeddings.
- **Smart Formatting:**
 - **Bold Keywords:** Highlights key concepts.
 - **Word Equations:** Formats chemical/biological processes clearly on separate lines.
 - **Clean Citations:** Removes inline clutter and provides a consolidated "Source" footer.
- **Hybrid Ingestion:** Handles both standard digital PDFs and scanned images using OCR (Tesseract).
- **Performance Benchmarking:** Integrated tools to measure latency and generation speed.

2. Technical Architecture & Component Breakdown

A. Tech Stack

- **Language:** Python 3.10+
- **UI Framework:** Streamlit
- **LLM Orchestration:** LangChain
- **Vector Database:** FAISS (Facebook AI Similarity Search) - CPU optimized.
- **Embeddings:** sentence-transformers/paraphrase-multilingual-MiniLM-L12-v2 (HuggingFace).
- **LLM Inference:** llama-cpp-python running a Quantized (GGUF) **Mistral-7B-Instruct** model.
- **OCR Engine:** Tesseract OCR (via pytesseract and pdf2image).

B. How It Works (The Pipeline)

1. Data Ingestion (`src/ingestion.py`):

- **Loading:** Recursively scans `data/raw/` for PDF files.
- **OCR processing:** If a PDF is scanned (text not extractable), it converts pages to images and uses **Tesseract** to extract text.
- **Chunking:** Splits text into 500-character chunks with 50-character overlap using `RecursiveCharacterTextSplitter`.
- **Vectorization:** Converts chunks into numerical vectors using the Embedding Model and saves them to a local FAISS index (`data/vectorized/`).

2. Retrieval (`src/retrieval.py`):

- When a user asks a question, the system converts the question into a vector.
- It searches the FAISS index for the top `k` (default 5) most similar chunks of text from the textbooks.

3. Generation (`src/generation.py`):

- **Prompt Engineering:** A strictly engineered prompt instructs the LLM to format the answer (bolding, spacing, equations) and ignore inline citations.
- **Inference:** The context and question are sent to the local Mistral 7B model.
- **Post-Processing:** A Regex cleaner aggressively strips any hallucinated inline citations from the text.
- **Footer Generation:** The system programmatically builds a "Source" list based on the metadata of the retrieved chunks.

4. UI (`app.py`):

- Provides a chat interface.
- Manages session state (chat history).
- Displays the final formatted answer and source citations.

3. Installation & Setup

Prerequisites

1. **Python 3.10+**
2. **Tesseract OCR:** Must be installed and added to System PATH.
3. **Poppler:** Required for processing PDF images (OCR).
4. **C++ Build Tools:** Required for compiling `llama-cpp-python` (Visual Studio Build Tools on Windows).

Installation Commands

Run these in your terminal within the project directory:

```
# 1. Install Python Dependencies  
pip install -r requirements.txt  
  
# 2. Download the Model  
# (Ensure 'mistral-7b-instruct-v0.1.Q4_K_M.gguf' is placed inside 'models/' directory)
```

4. How to Run

Step 1: Ingest Data (Prepare the Brain)

If you add new PDFs to `data/raw/`, run this command to update the database:

```
python src/ingestion.py
```

Output: This will create `index.faiss` and `index.pkl` in `data/vectorized/`.

Step 2: Launch the App

To start the Chat Interface:

```
streamlit run app.py
```

Access the app at: <http://localhost:8501>

Step 3: Run Benchmarks (Optional)

To test system performance (latency/speed) across 50 questions:

```
python benchmark_50.py
```

Results will be saved to `benchmark_50_results.csv`.

Step 4: Stopping the Application

To stop the application or any running script:

1. Click inside the terminal window where the app is running.
2. Press **Ctrl + C** on your keyboard.
3. The process will terminate, and you will see the command prompt again.

5. File Structure

```
Project_Root/  
|__ data/  
|   |__ raw/          # Place your NCERT PDFs here  
|   \__ vectorized/  # Generated FAISS index files  
|__ models/          # Place GGUF LLM models here  
|__ src/  
|   |__ ingestion.py # ETL Pipeline (PDF -> Vector DB)  
|   |__ retrieval.py # Search Logic  
|   |__ generation.py # LLM & Formatting Logic  
|   |__ pipeline.py  # Orchestrator  
|   |__ utils.py     # Helpers (Language detection)  
|   __ app.py        # Main Streamlit Application  
|   __ config.py    # Configuration (Paths, Prompts, Constants)  
|   __ requirements.txt # Python dependencies  
|   __ benchmark_50.py # Automated Stress Test  
|   __ PROJECT_REPORT.md # This Documentation
```